

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
```

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
import seaborn as sns
import matplotlib.pyplot as plt
```

```
df = pd.read_csv('/content/T_ML_Dataset151025.csv')
```

```
/tmp/ipython-input-4092072916.py:1: DtypeWarning: Columns (15) have mixed types. Specify dtype option on import or set low_memory=False.
 df = pd.read_csv('/content/T_ML_Dataset151025.csv')
```

```
# Show size and columns
print(df.shape)
df.head()
```

```
(594687, 28)
device_id      ts    power_w  voltage_v  current_a  relay_on  energy_wh_acc  firmware  location  illuminance_lux ...  outlet_id  is_110v  is_220v  Bracker_amp  max_watts  hour  weekday  Month  is_weekend  is_anomaly
0   1 2024-10-07 23:50:00  732.770186  118.880771  6.163908     1  6660455  1.12.3  77327        0.0  ...       1   1   0     30  3800  23      0   10      0      1
1   1 2024-10-07 23:50:00  768.632021  120.267027  6.391045     1  6658628  1.12.3  77327        0.0  ...       1   1   0     30  3800  23      0   10      0      1
2   1 2024-10-08 00:00:00  835.137039  119.387939  6.995154     1      139  1.12.3  77327        0.0  ...       1   1   0     30  3800  0      1   10      0      1
3   1 2024-10-08 00:10:00  619.227933  119.496472  5.181977     0      242  1.12.3  77327        0.0  ...       1   1   0     30  3800  0      1   10      0      0
4   1 2024-10-08 00:20:00  524.165806  118.527716  4.422306     1      330  1.12.3  77327        0.0  ...       1   1   0     30  3800  0      1   10      0      1
```

```
5 rows × 28 columns
```

```
df.info()
df.describe()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 594687 entries, 0 to 594686
Data columns (total 28 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   device_id        594687 non-null   int64  
 1   ts                594687 non-null   object  
 2   power_w          594687 non-null   float64 
 3   voltage_v         594687 non-null   float64 
 4   current_a         594687 non-null   float64 
 5   relay_on          594687 non-null   int64  
 6   energy_wh_acc    594687 non-null   int64  
 7   firmware          594687 non-null   object  
 8   location          594687 non-null   int64  
 9   illuminance_lux  594687 non-null   float64 
 10  presence          594687 non-null   int64  
 11  presence_confidence 594687 non-null   float64 
 12  temp_c_avg        594687 non-null   float64 
 13  Hostname          594687 non-null   object  
 14  vendor             594687 non-null   object  
 15  model              594687 non-null   object  
 16  device_type        594687 non-null   int64  
 17  area               594687 non-null   int64  
 18  outlet_id          594687 non-null   int64  
 19  is_110v            594687 non-null   int64  
 20  is_220v            594687 non-null   int64  
 21  Bracker_amp       594687 non-null   int64  
 22  max_watts          594687 non-null   int64  
 23  hour               594687 non-null   int64  
 24  weekday            594687 non-null   int64  
 25  Month              594687 non-null   int64  
 26  is_weekend         594687 non-null   int64  
 27  is_anomaly          594687 non-null   int64  
dtypes: float64(6), int64(17), object(5)
memory usage: 127.4+ MB

```

	device_id	power_w	voltage_v	current_a	relay_on	energy_wh_acc	location	illuminance_lux	presence	presence_confidence	...	outlet_id	is_110v	is_220v	Bracker_amp	max_watts	hour	weekday	Month	
count	594687.000000	594687.000000	594687.000000	594687.000000	594687.000000	5.946870e+05	594687.0	594687.000000	594687.000000	594687.000000	...	594687.000000	594687.000000	594687.000000	594687.000000	594687.000000	594687.000000	594687.000000	594687.000000	594687.000000
mean	2.772495	755.884727	119.998537	6.300036	0.849776	3.246057e+06	77327.0	177.830669	0.226075	0.435731	...	4.544991	0.291002	0.708998	37.089982	7344.990894	11.497964	2.996775	6.626491	
std	1.475042	154.625994	1.443347	1.291154	0.357291	1.946336e+06	0.0	257.105611	0.418289	0.247317	...	2.347872	0.454225	0.454225	9.567233	2930.795056	6.922330	2.001259	3.465153	
min	1.000000	488.753063	117.500005	3.991846	0.000000	8.200000e+01	77327.0	0.000000	0.000000	0.200000	...	1.000000	0.000000	0.000000	30.000000	3800.000000	0.000000	0.000000	1.000000	
25%	1.000000	638.631636	118.748183	5.319629	1.000000	1.529204e+06	77327.0	0.000000	0.000000	0.258118	...	1.000000	0.000000	0.000000	30.000000	3800.000000	5.000000	1.000000	4.000000	
50%	3.000000	745.410975	119.997166	6.212132	1.000000	3.232541e+06	77327.0	0.000000	0.000000	0.329961	...	6.000000	0.000000	1.000000	30.000000	6600.000000	11.000000	3.000000	7.000000	
75%	4.000000	850.103272	121.249308	7.084141	1.000000	4.935722e+06	77327.0	387.113497	0.000000	0.439590	...	6.000000	1.000000	1.000000	50.000000	11000.000000	17.000000	5.000000	10.000000	
max	5.000000	1178.743914	122.499999	10.023153	1.000000	6.667095e+06	77327.0	719.948380	1.000000	0.999999	...	7.000000	1.000000	1.000000	50.000000	11000.000000	23.000000	6.000000	12.000000	

8 rows × 23 columns

```
# Number of Nulls
df.isnull().sum()
```

```
0
device_id      0
ts              0
power_w        0
voltage_v      0
current_a      0
relay_on        0
energy_wh_acc  0
firmware        0
location         0
illuminance_lux 0
presence         0
presence_confidence 0
temp_c_avg      0
Hostname        0
vendor          0
model           0
device_type     0
area            0
outlet_id       0
is_110v         0
is_220v         0
Bracker_amp    0
max_watts       0
hour            0
weekday         0
Month           0
is_weekend      0
is_anomaly      0
```

dtype: int64

```
df.dropna(inplace=True)
print("Filas con valores nulos eliminadas.")
```

Filas con valores nulos eliminadas.

```
# =====
# [1] Análisis de balance de clases
# =====
import matplotlib.pyplot as plt
import seaborn as sns

# Contar los valores
conteo = df['is_anomaly'].value_counts().sort_index()

# Mostrar resultados numéricos
print("Class distribution (is_anomaly):")
for etiqueta, cantidad in conteo.items():
    print(f" Class {etiqueta}: {cantidad} rows ({cantidad/len(df)*100:.2f}%)")

# =====
# [2] Visualización con gráfico de barras
# =====
plt.figure(figsize=(6,4))
sns.barplot(x=conteo.index, y=conteo.values, palette=['#4CAF50','#F44336'])
plt.title('Class Distribution: Normal (0) vs Abnormal (1)', fontsize=14)
plt.xlabel('Clase (is_anomaly)', fontsize=12)
plt.ylabel('Número de Filas', fontsize=12)
plt.xticks([0,1], ['Normal (0)', 'Abnormal (1)'])
for i, v in enumerate(conteo.values):
    plt.text(i, v + len(df)*0.005, str(v), ha='center', fontsize=10)
plt.show()

# =====
# [3] Evaluación de balance
# =====
minor = conteo.min()
mayor = conteo.max()
ratio = minor / mayor
```

```

if ratio < 0.5:
    print(f"⚠️ Unbalanced dataset: minor class = {minor}, Senior class = {major} (ratio = {ratio:.2f})")
    print("👉 Sugerencia: usar SMOTE, RandomUnderSampler o class_weight='balanced' Model.")
else:
    print(f"✅ Dataset balanceado (ratio = {ratio:.2f})")

```


Class distribution (is_anomaly):
 Class 0: 126316 rows (21.24%)
 Class 1: 468371 rows (78.76%)
/tmp/ipython-input-1737811432.py:19: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

sns.barplot(x=conteo.index, y=conteo.values, palette=['#4CAF50', '#F44336'])
/usr/local/lib/python3.12/dist-packages/IPython/core/pylabtools.py:151: UserWarning: Glyph 128994 (\N{LARGE GREEN CIRCLE}) missing from font(s) DejaVu Sans.
fig.canvas.print_figure(bytes_io, **kw)
/usr/local/lib/python3.12/dist-packages/IPython/core/pylabtools.py:151: UserWarning: Glyph 128308 (\N{LARGE RED CIRCLE}) missing from font(s) DejaVu Sans.
fig.canvas.print_figure(bytes_io, **kw)

```



⚠️ Unbalanced dataset: minor class = 126316, Senior class = 468371 (ratio = 0.27)
👉 Sugerencia: usar SMOTE, RandomUnderSampler o class_weight='balanced' Model.

```

# =====
# [1] Instalar e importar SMOTE
# =====
!pip install imbalanced-learn --quiet
from imblearn.over_sampling import SMOTE
import pandas as pd

# =====
# [2] Aplicar balanceo con SMOTE
# =====
smote = SMOTE(random_state=42)
X_bal, y_bal = smote.fit_resample(X, y)

# =====
# [3] Reconstruir el DataFrame balanceado
# =====
# X_bal es un numpy array, y_bal es un vector → hay que unirlos
df_bal = pd.DataFrame(X_bal, columns=X.columns)
df_bal['is_anomaly'] = y_bal

# =====
# [4] Verificar el nuevo balance
# =====
conteo = df_bal['is_anomaly'].value_counts()
print("🆕 Nueva distribución tras SMOTE:")
for k, v in conteo.items():
    print(f" Clase {k}: {v} registros ({v/len(df_bal)*100:.2f}%)")


# =====
# [5] Dividir dataset balanceado
# =====
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(
    df_bal.drop('is_anomaly', axis=1),
    df_bal['is_anomaly'],
    test_size=0.2,
    random_state=42,
    stratify=df_bal['is_anomaly']
)

print(f"\n✅ Dataset listo. Tamaño final:")
print(f" Train: {len(X_train)} | Test: {len(X_test)}")

```

🆕 Nueva distribución tras SMOTE:
 Clase 1: 468371 registros (50.00%)
 Clase 0: 468371 registros (50.00%)

Dataset listo. Tamaño final:
Train: 749393 | Test: 187349

```
# =====
# [1] Análisis de balance de clases
# =====
import matplotlib.pyplot as plt
import seaborn as sns

# Contar los valores
conteo = df_bal['is_anomaly'].value_counts().sort_index()

# Mostrar resultados numéricos
print("📊 Class distribution (is_anomaly):")
for etiqueta, cantidad in conteo.items():
    print(f"  Class {etiqueta}: {cantidad} rows ({cantidad/len(df_bal)*100:.2f}%)")

# =====
# [2] Visualización con gráfico de barras
# =====
plt.figure(figsize=(6,4))
sns.barplot(x=conteo.index, y=conteo.values, palette=['#4CAF50','#F44336'])
plt.title('Class Distribution: Normal (0) vs Abnormal (1)', fontsize=14)
plt.xlabel('Clase (is_anomaly)', fontsize=12)
plt.ylabel('Número de Filas', fontsize=12)
plt.xticks([0,1], ['🟡 Normal (0)', '🔴 Abnormal (1)'])
for i, v in enumerate(conteo.values):
    plt.text(i, v + len(df)*0.005, str(v), ha='center', fontsize=10)
plt.show()


# =====
# [3] Evaluación de balance
# =====
minor = conteo.min()
major = conteo.max()
ratio = minor / major

if ratio < 0.5:
    print("⚠️ Unbalanced dataset: minor class = {minor}, Senior class = {major} (ratio = {ratio:.2f})")
    print("👉 Sugerencia: usar SMOTE, RandomUnderSampler o class_weight='balanced' Model.")
else:
    print("✅ Dataset balanceado (ratio = {ratio:.2f})")
```

📊 Class distribution (is_anomaly):
Class 0: 468371 rows (50.00%)
Class 1: 468371 rows (50.00%)
/tmp/ipython-input-2088039122.py:19: FutureWarning:


Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=conteo.index, y=conteo.values, palette=['#4CAF50','#F44336'])
/usr/local/lib/python3.12/dist-packages/IPython/core/pylabtools.py:151: UserWarning: Glyph 128994 (\N{LARGE GREEN CIRCLE}) missing from font(s) DejaVu Sans.
fig.canvas.print_figure(bytes_io, **kw)
/usr/local/lib/python3.12/dist-packages/IPython/core/pylabtools.py:151: UserWarning: Glyph 128308 (\N{LARGE RED CIRCLE}) missing from font(s) DejaVu Sans.
fig.canvas.print_figure(bytes_io, **kw)
```




Dataset balanceado (ratio = 1.00)


```
sns.histplot(df_bal['power_w'], bins=50, kde=True)
plt.show()
```



```
sns.histplot(df_bal['voltage_v'], bins=50, kde=True)  
plt.show()
```




```
sns.histplot(df_bal['current_a'], bins=50, kde=True)  
plt.show()
```




```
# df_bal['ts'] = pd.to_datetime(df_bal['ts']) # This line is removed  
# df_bal['hour'] = df_bal['ts'].dt.hour # This line is removed
```

```
plt.figure(figsize=(12,6))  
sns.boxplot(x='hour', y='power_w', data=df_bal)  
plt.title('Distribution of consumption (power_w) by hour of the day')  
plt.show()
```




```
df_bal['is_daytime'] = ((df_bal['hour'] >= 7) & (df_bal['hour'] <= 19)).astype(int)
```


```
plt.figure(figsize=(8,5))
sns.boxplot(x='is_daytime', y='power_w', data=df_bal)
plt.title('Comparate of Consume (power_w): Day (1) vs Night (0)')
plt.xlabel('1 = Day, 0 = Night')
plt.show()
```



```
plt.figure(figsize=(10,6))
sns.scatterplot(x='illuminance_lux', y='power_w', hue='is_daytime', data=df_bal)
plt.title('Relación entre iluminancia y consumo de potencia')
plt.show()
```



```
context_vars = ['hour', 'is_daytime', 'illuminance_lux', 'power_w', 'voltage_v', 'current_a']
sns.heatmap(df_bal[context_vars].corr(), annot=True, cmap='coolwarm')
plt.title('Correlación entre hora, luz y variables eléctricas')
plt.show()
```



```
# Clean Data
```

```
# Revisar columnas actuales
print("Columnas originales:", df_bal.columns.tolist())

# Eliminar columnas irrelevantes
columns_to_drop = ['firmware', 'vendor', 'Hostname', 'location']

# Eliminar una de is_110v o is_220v (mantener solo una)
# Ejemplo: creamos una columna combinada y borramos ambas
df_bal['voltage_type'] = df['is_220v'].apply(lambda x: '220V' if x == 1 else '110V')

# Luego eliminamos las originales
columns_to_drop.extend(['is_110v', 'is_220v'])

# Aplicar eliminación
df_bal.drop(columns=columns_to_drop, inplace=True, errors='ignore')

print("Columnas después de limpieza:", df_bal.columns.tolist())

Columnas originales: ['power_w', 'voltage_v', 'current_a', 'relay_on', 'energy_wh_acc', 'illuminance_lux', 'presence', 'presence_confidence', 'temp_c_avg', 'device_type', 'area', 'outlet_id', 'is_110v', 'is_220v', 'Bracker_amp', 'max_watts', 'hour', 'weekday', 'Month', 'is_week'
Columnas después de limpieza: ['power_w', 'voltage_v', 'current_a', 'relay_on', 'energy_wh_acc', 'illuminance_lux', 'presence', 'presence_confidence', 'temp_c_avg', 'device_type', 'area', 'outlet_id', 'Bracker_amp', 'max_watts', 'hour', 'weekday', 'Month', 'is_weekend', 'is_and']
```

```
# Aplicar One-Hot Encoding a 'device_type' y 'area'
df_encoded = pd.get_dummies(df_bal, columns=['device_type', 'area'], prefix=['dev', 'area'])

# Verificar las nuevas columnas
print("Nuevas columnas:", [col for col in df_encoded.columns if 'dev_' in col or 'area_' in col])

Nuevas columnas: ['dev_1', 'dev_3', 'dev_4', 'dev_5', 'dev_6', 'dev_7', 'dev_8', 'area_1', 'area_2', 'area_3', 'area_4']
```

```
print(df_encoded.head())
```

```
power_w    voltage_v   current_a   relay_on   energy_wh_acc \
0  732.770186  118.880771  6.163908      1     6660455
1  768.632021  120.267027  6.391045      1     6658628
2  835.137039  119.387939  6.995154      1     139
3  619.227933  119.496472  5.181977      0     242
4  524.165806  118.527716  4.422306      1     330

illuminance_lux  presence  presence_confidence  temp_c_avg  outlet_id ...
0       0.0        0          0.276138  16.479134      1 ...
1       0.0        0          0.375432  16.308392      1 ...
2       0.0        0          0.379196  14.872813      1 ...
3       0.0        0          0.225525  14.872813      1 ...
4       0.0        0          0.304784  14.872813      1 ...

dev_3  dev_4  dev_5  dev_6  dev_7  dev_8  area_1  area_2  area_3  area_4
0  False  False  False  False  False  False  True  False  False  False
1  False  False  False  False  False  False  False  True  False  False
2  False  False  False  False  False  False  True  False  False  False
3  False  False  False  False  False  False  True  False  False  False
4  False  False  False  False  False  False  True  False  False  False
```

[5 rows x 30 columns]

```
print(f"Dataset cargado: {df_bal.shape[0]} registros, {df_bal.shape[1]} columnas")
```

```
# ✅ Seleccionar features y variable objetivo
```

```
# -----
# Usamos variables eléctricas y contextuales como features,
# y 'is_anomaly' como etiqueta (target)
target = 'is_anomaly'
```

```
# Excluir columnas no predictivas o de identificación
```

```
excluded_cols = ['device_id', 'ts']
X = df.drop(columns=excluded_cols + [target], errors='ignore')
y = df[target]
```

```
print(f"Variables predictoras: {X.shape[1]} | Target: '{target}'")
```

```
Dataset cargado: 936742 registros, 21 columnas
Variables predictoras: 25 | Target: 'is_anomaly'
```

```
if 'voltage_type' in X.columns:
    X['voltage_type'] = X['voltage_type'].replace({'110V': 0, '220V': 1})
```

```
# Detectar columnas no numéricas
non_numeric_cols = X.select_dtypes(exclude=[np.number]).columns.tolist()
print("Columnas no numéricas eliminadas:", non_numeric_cols)
```

```
# Eliminar columnas de texto
X = X.drop(columns=non_numeric_cols, errors='ignore')
```

```
Columnas no numéricas eliminadas: ['firmware', 'Hostname', 'vendor', 'model']
```

```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
from sklearn.model_selection import train_test_split
```

```
# División del dataset: 80% entrenamiento, 20% prueba
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled,
    y,
    test_size=0.2,      # 20% de los datos para prueba
    random_state=42,    # reproducibilidad
    stratify=y          # mantiene la proporción de anomalías
)
```

```
print(f"Entrenamiento: {X_train.shape[0]} muestras | Prueba: {X_test.shape[0]} muestras")
```

Entrenamiento: 475749 muestras | Prueba: 118938 muestras

```
from sklearn.ensemble import RandomForestClassifier
```

```

model = RandomForestClassifier(
    n_estimators=200,
    random_state=42,
    class_weight='balanced'
)

model.fit(X_train, y_train)
print("✅ Modelo entrenado correctamente.")

```

✅ Modelo entrenado correctamente.

```

from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
import seaborn as sns
import matplotlib.pyplot as plt

y_pred = model.predict(X_test)

# Métricas principales
print("🎯 Accuracy: {:.4f}".format(accuracy_score(y_test, y_pred)))
print("\n📊 Classification Report:")
print(classification_report(y_test, y_pred))


# Matriz de confusión
plt.figure(figsize=(6,5))
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, fmt='d', cmap='Blues')
plt.title('Matriz de Confusión - Detección de Anomalías')
plt.xlabel('Predicción')
plt.ylabel('Real')
plt.show()

```

🎯 Accuracy: 1.0000

	precision	recall	f1-score	support
0	1.00	1.00	1.00	25263
1	1.00	1.00	1.00	93675
accuracy			1.00	118938
macro avg	1.00	1.00	1.00	118938
weighted avg	1.00	1.00	1.00	118938

Matriz de Confusión - Detección de Anomalías



Start coding or [generate](#) with AI.

```

# Ejemplo: registro de HVAC
nuevo_registro = {
    'power_w': 0,
    'voltage_v': 0,
    'current_a': 0,
    'relay_on': 1,
    'energy_wh_acc': 128200,
    'presence': 1,
    'presence_confidence': 0.9,
    'temp_c_avg': 23.5,
    'Bracker_amp': 30,
    'max_watts': 6000,
    'power_rate': 0 / 6000,  # 0.8
    'hour': 14,
    'is_daytime': 1,
    'weekday': 2,
    'area_1': 0, 'area_2': 1, 'area_3': 0,  # depende del encoding
    'device_type_1': 0, 'device_type_2': 1, 'device_type_3': 0,
}

```

```

'device_type_4': 0, 'device_type_5': 0, 'device_type_6': 0,
'device_type_7': 0, 'device_type_8': 0,
'voltage_type': 1 # 220V
}

# Nuevo registro (dentro de rango normal)
nuevo_registro = {
    'power_w': 4200, # normal
    'voltage_v': 230, # normal (200-240 V)
    'current_a': 18.3, # normal (15-30 A)
    'relay_on': 1,
    'energy_wh_acc': 550000,
    'presence': 1,
    'presence_confidence': 0.95,
    'temp_c_avg': 45.2,
    'Bracker_amp': 30,
    'max_watts': 5500,
    'power_rate': 4200 / 5500, # =0.76
    'hour': 10,
    'is_daytime': 1,
    'weekday': 2,
    'voltage_type': 1, # 220 V
    # ◇ Encoding de ejemplo
    'area_1': 1, 'area_2': 0, 'area_3': 0,
    'device_type_1': 1, 'device_type_2': 0, 'device_type_3': 0,
    'device_type_4': 0, 'device_type_5': 0, 'device_type_6': 0,
    'device_type_7': 0, 'device_type_8': 0
}

```

```

df_new = pd.DataFrame([nuevo_registro])

# Ensure the new data frame has the same columns as X, in the same order
# This is crucial because the scaler was fitted on X
missing_cols = set(X.columns) - set(df_new.columns)
for c in missing_cols:
    df_new[c] = 0 # Add missing columns with a default value (e.g., 0 for one-hot encoded)

# Reorder columns to match the order of X
df_new = df_new[X.columns]

# Escalar con el mismo StandardScaler
df_scaled = scaler.transform(df_new)

```

```

# Predicción
prediccion = model.predict(df_scaled)
probabilidad = model.predict_proba(df_scaled)

# Interpretar el resultado
estado = "🔴 Normal (0)" if int(prediccion[0]) == 0 else "🔴 Abnormal (1)"

print(f"Predicción: {estado}")
print(f"Probabilidad de anomalía: {probabilidad[0][1]*100:.2f}%")

```

```

Predicción: 🔴 Abnormal (1)
Probabilidad de anomalía: 98.00%

```

```
!pip install lazypredict
```

```

Requirement already satisfied: lazypredict in /usr/local/lib/python3.12/dist-packages (0.2.16)
Requirement already satisfied: click in /usr/local/lib/python3.12/dist-packages (from lazypredict) (8.3.0)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.12/dist-packages (from lazypredict) (1.6.1)
Requirement already satisfied: pandas in /usr/local/lib/python3.12/dist-packages (from lazypredict) (2.2.2)
Requirement already satisfied: tqdm in /usr/local/lib/python3.12/dist-packages (from lazypredict) (4.67.1)
Requirement already satisfied: joblib in /usr/local/lib/python3.12/dist-packages (from lazypredict) (1.5.2)
Requirement already satisfied: lightgbm in /usr/local/lib/python3.12/dist-packages (from lazypredict) (4.6.0)
Requirement already satisfied: xgboost in /usr/local/lib/python3.12/dist-packages (from lazypredict) (3.0.5)
Requirement already satisfied: pytest-runner in /usr/local/lib/python3.12/dist-packages (from lazypredict) (6.0.1)
Requirement already satisfied: mlflow>=2.0.0 in /usr/local/lib/python3.12/dist-packages (from lazypredict) (3.5.0)
Requirement already satisfied: mlflow-skiminy==3.5.0 in /usr/local/lib/python3.12/dist-packages (from mlflow>=2.0.0->lazypredict) (3.5.0)
Requirement already satisfied: mlflow-tracing>=3.5.0 in /usr/local/lib/python3.12/dist-packages (from mlflow>=2.0.0->lazypredict) (3.5.0)
Requirement already satisfied: Flask-CORS<7 in /usr/local/lib/python3.12/dist-packages (from mlflow>=2.0.0->lazypredict) (6.0.1)
Requirement already satisfied: Flask4<4 in /usr/local/lib/python3.12/dist-packages (from mlflow>=2.0.0->lazypredict) (3.1.2)
Requirement already satisfied: alembic!=1.10.0,<2 in /usr/local/lib/python3.12/dist-packages (from mlflow>=2.0.0->lazypredict) (1.16.5)
Requirement already satisfied: cryptography<47,>=43.0.0 in /usr/local/lib/python3.12/dist-packages (from mlflow>=2.0.0->lazypredict) (43.0.3)
Requirement already satisfied: docker<8,>=4.0.0 in /usr/local/lib/python3.12/dist-packages (from mlflow>=2.0.0->lazypredict) (7.1.0)
Requirement already satisfied: fastmpc3,>=2.0.0 in /usr/local/lib/python3.12/dist-packages (from mlflow>=2.0.0->lazypredict) (2.12.4)
Requirement already satisfied: graphene4 in /usr/local/lib/python3.12/dist-packages (from mlflow>=2.0.0->lazypredict) (3.4.3)
Requirement already satisfied: gunicorn<24 in /usr/local/lib/python3.12/dist-packages (from mlflow>=2.0.0->lazypredict) (23.0.0)
Requirement already satisfied: matplotlib4 in /usr/local/lib/python3.12/dist-packages (from mlflow>=2.0.0->lazypredict) (3.10.0)
Requirement already satisfied: numpy<3 in /usr/local/lib/python3.12/dist-packages (from mlflow>=2.0.0->lazypredict) (2.0.2)
Requirement already satisfied: pyarrow<22,>=4.0.0 in /usr/local/lib/python3.12/dist-packages (from mlflow>=2.0.0->lazypredict) (18.1.0)
Requirement already satisfied: scipy<2 in /usr/local/lib/python3.12/dist-packages (from mlflow>=2.0.0->lazypredict) (1.16.2)
Requirement already satisfied: sqlalchemy<3,>=1.4.0 in /usr/local/lib/python3.12/dist-packages (from mlflow>=2.0.0->lazypredict) (2.0.43)
Requirement already satisfied: cachetools<7,>=5.0.0 in /usr/local/lib/python3.12/dist-packages (from mlflow>=2.0.0->lazypredict) (5.5.2)
Requirement already satisfied: cloudpickle<4 in /usr/local/lib/python3.12/dist-packages (from mlflow-skiminy==3.5.0->mlflow>=2.0.0->lazypredict) (3.1.1)
Requirement already satisfied: databricks-sd<1,>=0.20.0 in /usr/local/lib/python3.12/dist-packages (from mlflow-skiminy==3.5.0->mlflow>=2.0.0->lazypredict) (0.68.0)
Requirement already satisfied: fastapi<1 in /usr/local/lib/python3.12/dist-packages (from mlflow-skiminy==3.5.0->mlflow>=2.0.0->lazypredict) (0.118.2)
Requirement already satisfied: gitpython<4,>=3.1.9 in /usr/local/lib/python3.12/dist-packages (from mlflow-skiminy==3.5.0->mlflow>=2.0.0->lazypredict) (3.1.45)

```

```

Requirement already satisfied: importlib_machinery<4.7.0,>=3.7.0 in /usr/local/lib/python3.12/dist-packages (from mlflow-skinny==3.5.0->mlflow>=2.0.0->lazypredict) (8.7.0)
Requirement already satisfied: opentelemetry-api<3.2,>=1.9.0 in /usr/local/lib/python3.12/dist-packages (from mlflow-skinny==3.5.0->mlflow>=2.0.0->lazypredict) (1.37.0)
Requirement already satisfied: opentelemetryproto<3.2,>=1.9.0 in /usr/local/lib/python3.12/dist-packages (from mlflow-skinny==3.5.0->mlflow>=2.0.0->lazypredict) (1.38.0)
Requirement already satisfied: opentelemetry-sdk<3.2,>=1.9.0 in /usr/local/lib/python3.12/dist-packages (from mlflow-skinny==3.5.0->mlflow>=2.0.0->lazypredict) (1.37.0)
Requirement already satisfied: packaging<2.0 in /usr/local/lib/python3.12/dist-packages (from mlflow-skinny==3.5.0->mlflow>=2.0.0->lazypredict) (25.0)
Requirement already satisfied: protobuf<7,>=3.12.0 in /usr/local/lib/python3.12/dist-packages (from mlflow-skinny==3.5.0->mlflow>=2.0.0->lazypredict) (5.29.5)
Requirement already satisfied: pydantic<3,>=1.10.8 in /usr/local/lib/python3.12/dist-packages (from mlflow-skinny==3.5.0->mlflow>=2.0.0->lazypredict) (2.11.10)
Requirement already satisfied: python-donetv<2,>=0.19.0 in /usr/local/lib/python3.12/dist-packages (from mlflow-skinny==3.5.0->mlflow>=2.0.0->lazypredict) (1.1.1)
Requirement already satisfied: pyyaml<7,>=5.1 in /usr/local/lib/python3.12/dist-packages (from mlflow-skinny==3.5.0->mlflow>=2.0.0->lazypredict) (6.0.3)
Requirement already satisfied: requests<3,>=2.17.3 in /usr/local/lib/python3.12/dist-packages (from mlflow-skinny==3.5.0->mlflow>=2.0.0->lazypredict) (2.32.4)
Requirement already satisfied: sqlparse<1,>>0.4.0 in /usr/local/lib/python3.12/dist-packages (from mlflow-skinny==3.5.0->mlflow>=2.0.0->lazypredict) (0.5.3)
Requirement already satisfied: typing_extensions<5,>=4.0.0 in /usr/local/lib/python3.12/dist-packages (from mlflow-skinny==3.5.0->mlflow>=2.0.0->lazypredict) (4.15.0)
Requirement already satisfied: uvicorn<1 in /usr/local/lib/python3.12/dist-packages (from mlflow-skinny==3.5.0->mlflow>=2.0.0->lazypredict) (0.37.0)
Requirement already satisfied: python-dateutil<2.8.2 in /usr/local/lib/python3.12/dist-packages (from pandas->lazypredict) (2.9.0.post0)
Requirement already satisfied: pytz<2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas->lazypredict) (2025.2)
Requirement already satisfied: tzdata<2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas->lazypredict) (2025.2)
Requirement already satisfied: threadpoolctl<=3.1.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn->lazypredict) (3.6.0)
Requirement already satisfied: nvidia-nccl-cu12 in /usr/local/lib/python3.12/dist-packages (from xgboost->lazypredict) (2.27.3)
Requirement already satisfied: Make in /usr/local/lib/python3.12/dist-packages (from alembic<1.10.0,>2->mlflow>=2.0.0->lazypredict) (1.3.10)
Requirement already satisfied: cffi<1.12 in /usr/local/lib/python3.12/dist-packages (from cryptography<47,>>43.0.0->mlflow>=2.0.0->lazypredict) (2.0.0)
Requirement already satisfied: urllib3<1.26.0 in /usr/local/lib/python3.12/dist-packages (from docker<8,>>4.0.0->mlflow>=2.0.0->lazypredict) (2.5.0)
Requirement already satisfied: authlib<1.5.2 in /usr/local/lib/python3.12/dist-packages (from fastmpc<3,>>2.0.0->mlflow>=2.0.0->lazypredict) (1.6.5)
Requirement already satisfied: cyclops<=3.0.0 in /usr/local/lib/python3.12/dist-packages (from fastmpc<3,>>2.0.0->mlflow>=2.0.0->lazypredict) (3.24.0)
Requirement already satisfied: exceptiongroup<=1.2.2 in /usr/local/lib/python3.12/dist-packages (from fastmpc<3,>>2.0.0->mlflow>=2.0.0->lazypredict) (1.3.0)
Requirement already satisfied: httpx<0.28.1 in /usr/local/lib/python3.12/dist-packages (from fastmpc<3,>>2.0.0->mlflow>=2.0.0->lazypredict) (0.28.1)
Requirement already satisfied: mcp<2.0.0,>>1.12.4 in /usr/local/lib/python3.12/dist-packages (from fastmpc<3,>>2.0.0->mlflow>=2.0.0->lazypredict) (1.16.0)
Requirement already satisfied: openapi-core<0.19.5 in /usr/local/lib/python3.12/dist-packages (from fastmpc<3,>>2.0.0->mlflow>=2.0.0->lazypredict) (0.19.5)
Requirement already satisfied: onenano-nvdantic<=0.5.1 in /usr/local/lib/python3.12/dist-packages (from fastmpc<3,>>2.0.0->mlflow>=2.0.0->lazypredict) (0.5.1)

```

```

import gc, numpy as np, pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, f1_score, recall_score, precision_score
from sklearn.feature_selection import mutual_info_classif
from sklearn.linear_model import LogisticRegression, SGDClassifier
from sklearn.ensemble import RandomForestClassifier, ExtraTreesClassifier, GradientBoostingClassifier
from sklearn.svm import LinearSVC

# 1) Carga
df = pd.read_csv('/content/T_ML_Dataset151025.csv')

assert 'is_anomaly' in df.columns, "Falta is_anomaly"
y = df['is_anomaly'].astype(int)

# Quita columnas no predictivas / textuales
drop_cols = [c for c in ['device_id', 'ts', 'Hostname', 'firmware', 'vendor', 'location', 'model'] if c in df.columns]
X = df.drop(columns=drop_cols + ['is_anomaly'], errors='ignore')

# Normaliza voltage_type si quedó como texto
if 'voltage_type' in X.columns and X['voltage_type'].dtype == 'object':
    X['voltage_type'] = X['voltage_type'].replace({'110V':0, '220V':1})

# Elimina no numéricas y NaN/Inf
X = X.select_dtypes(include=[np.number]).replace([np.inf, -np.inf], np.nan).fillna(0)

# 2) Muestreo estratificado: máx 15k por clase (ajusta si necesitas)
max_per_class = 15000
idx_keep = []
for c, idx in y.groupby(y).groups.items():
    take = np.random.RandomState(42).choice(idx, size=min(len(idx), max_per_class), replace=False)
    idx_keep.append(take)
idx_keep = np.concatenate(idx_keep)
X_s = X.iloc[idx_keep].astype('float32')
y_s = y.iloc[idx_keep].astype('int8')

print("Tamaño tras muestreo:", X_s.shape, dict(pd.Series(y_s).value_counts()))

# 3) Split antes del escalado
X_train, X_test, y_train, y_test = train_test_split(
    X_s, y_s, test_size=0.2, stratify=y_s, random_state=42
)

# 4) Selección de features (quedarse con las K mejores)
K = min(60, X_train.shape[1]) # hasta 60, ajusta según RAM
mi = mutual_info_classif(X_train, y_train, random_state=42)
topk_idx = np.argsort(mi)[-K:]
cols_k = X_train.columns[topk_idx]
X_train = X_train[cols_k].copy()
X_test = X_test[cols_k].copy()

# 5) Escalado
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# 6) Benchmark ligero de modelos
models = {
    "LogReg": LogisticRegression(max_iter=1000, class_weight='balanced', n_jobs=None),
    "SGD": SGDClassifier(loss='log_loss', class_weight='balanced', random_state=42),
    "RF": RandomForestClassifier(n_estimators=200, max_depth=16, class_weight='balanced', random_state=42),
    "ET": ExtraTreesClassifier(n_estimators=300, max_depth=16, random_state=42),
    "GB": GradientBoostingClassifier(random_state=42),
    "LinSVC": LinearSVC(dual='auto', random_state=42)
}

```

```

rows = []
for name, m in models.items():
    try:
        m.fit(X_train, y_train)
        y_pred = m.predict(X_test)
        rows.append({
            "Model": name,
            "Accuracy": accuracy_score(y_test, y_pred),
            "F1": f1_score(y_test, y_pred, zero_division=0),
            "Recall": recall_score(y_test, y_pred, zero_division=0),
            "Precision": precision_score(y_test, y_pred, zero_division=0)
        })
    except Exception as e:
        rows.append({"Model": name, "Accuracy": np.nan, "F1": np.nan, "Recall": np.nan, "Precision": np.nan})
        print(f"X {name} falló:", repr(e))
finally:
    gc.collect()

bench = pd.DataFrame(rows).sort_values('F1', ascending=False)
print(bench)
bench.to_csv('/content/model_benchmark_manual.csv', index=False)
print("Ranking guardado en /content/model_benchmark_manual.csv")

```

```

/tmp/ipython-input-3277236342.py:11: DtypeWarning: Columns (15) have mixed types. Specify dtype option on import or set low_memory=False.
df = pd.read_csv('/content/T_ML_Dataset151025.csv')
Tamaño tras muestreo: (30000, 20) {0: np.int64(15000), 1: np.int64(15000)}
   Model  Accuracy      F1     Recall  Precision
4      GB  1.000000  1.000000  1.000000
2      RF  1.000000  1.000000  1.000000
3      ET  0.993500  0.993457  0.987000  1.000000
0  LogReg  0.983000  0.982706  0.966000  1.000000
1    SGD  0.981667  0.981375  0.966000  0.997247
5  LinSVC  0.978167  0.977679  0.956333  1.000000
Ranking guardado en /content/model_benchmark_manual.csv

```

Model	Accuracy	F1	Recall	Precision
RF (RandomForest)	1.000	1.000	1.000	1.000
GB (GradientBoosting)	1.000	1.000	1.000	1.000
ET (ExtraTrees)	0.993	0.993	0.987	1.000
LogReg	0.983	0.983	0.966	1.000
SGD	0.982	0.981	0.966	0.997
LinSVC	0.978	0.978	0.956	1.000

```

from sklearn.model_selection import cross_val_score
for name, model in [("RF", RandomForestClassifier(n_estimators=200, max_depth=16, random_state=42)),
                     ("ET", ExtraTreesClassifier(n_estimators=300, max_depth=16, random_state=42)),
                     ("GB", GradientBoostingClassifier(random_state=42))]:
    scores = cross_val_score(model, X_train, y_train, cv=5, scoring='f1')
    print(f"{name}: mean F1={scores.mean():.3f} ± {scores.std():.3f}")

```

```

RF: mean F1=1.000 ± 0.000
ET: mean F1=0.993 ± 0.001
GB: mean F1=1.000 ± 0.000

```

```

import joblib
import pandas as pd
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split

# Assuming X and y are already defined and preprocessed from previous steps
# If not, you would need to include the data loading and preprocessing here
# For demonstration, let's use the sampled and scaled data from the previous benchmark cell
# You might want to adjust this based on which data you want to use for the final model

# Re-split the sampled data to ensure X_train, X_test, y_train, y_test are available

```